

Task 1: Sparse Networks from Scratch

In the first task, students will train sparse networks with a fixed random connectivity mask. The mask M is generated at initialization and remains unchanged throughout training. Students should compare at least two optimization methods: SGD with momentum and Adam. The goal is to investigate how different optimizers behave when training highly sparse networks. Suggested questions include: - How does sparsity affect convergence speed? - How does accuracy degrade as sparsity increases? - Do different optimizers behave differently under high sparsity?

```
# from google.colab import drive
# drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive")`

Load libraries

```
import os
import sys
import torch
import torch.nn as nn
import torch.optim as optim
import joblib
import torchvision
import torchvision.transforms as transforms

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# quarto preview 01_task_1.ipynb --to pdf
```

```
# model_path = "/content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project"
# sys.path.append(model_path)
```

```
model_path = "."
```

```
from models import ResNetCIFAR
```

```
use_pin_memory = torch.cuda.is_available()
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
num_workers = min(2, os.cpu_count())
```

Load data

```
# Standard CIFAR-10 transforms
```

```
transform = transforms.Compose(
```

```
[
```

```
    transforms.ToTensor(),
```

```
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
```

```
]
```

```
)
```

```
trainset = torchvision.datasets.CIFAR10(
```

```
    root=f"{model_path}/data", train=True, download=True, transform=transform
```

```
)
```

```
trainloader = torch.utils.data.DataLoader(
```

```
    trainset,
```

```
    batch_size=128,
```

```
    shuffle=True,
```

```
    num_workers=num_workers,
```

```
    pin_memory=use_pin_memory,
```

```
)
```

```
testset = torchvision.datasets.CIFAR10(
```

```
    root=f"{model_path}/data", train=False, download=True, transform=transform
```

```
)
```

```
testloader = torch.utils.data.DataLoader(
```

```
    testset,
```

```
    batch_size=128,
```

```

        shuffle=False,
        num_workers=num_workers,
        pin_memory=use_pin_memory,
    )

```

Configuration

```

sparsities = [0.90, 0.95, 0.98]
optimizers = ["sgd", "adam"]

```

Main Experiment

```

epochs = 15
criterion = nn.CrossEntropyLoss()

```

```

output_dir = "/content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project/results"

```

Helper Functions

```

def get_static_masks(model, sparsity):
    """Generates the fixed binary mask M"""
    masks = {}
    for name, param in model.named_parameters():
        if "conv" in name or "linear" in name:
            mask = torch.rand(param.size()).to(device)
            mask = (mask > sparsity).float()
            masks[name] = mask
        else:
            masks[name] = torch.ones(param.size()).to(device)
    return masks

@torch.no_grad()
def apply_masks(model, masks):
    """Applies W = M * theta"""
    for name, param in model.named_parameters():
        param.mul_(masks[name])

```

Actual experiment

This section runs the full experimental loop over sparsity levels and optimizers. For each configuration, a model is initialized, a fixed random mask is applied, and training proceeds while enforcing the same sparse connectivity throughout. Performance metrics are tracked at every epoch and saved for later analysis.

```
for s in sparsities:
    for opt_name in optimizers:
        print(f"\nStarting Experiment: Sparsity={s*100}%, Optimizer={opt_name}")

        # Initialize
        model = ResNetCIFAR().to(device)
        masks = get_static_masks(model, s)
        apply_masks(model, masks)

        if opt_name == "sgd":
            optimizer = optim.SGD(
                model.parameters(), lr=0.1, momentum=0.9, weight_decay=5e-4
            )
        else:
            optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=5e-4)

        history = {"train_loss": [], "train_acc": [], "test_loss": [], "test_acc": []}

        for epoch in range(epochs):
            # --- Training Phase ---
            model.train()
            train_loss, train_correct, train_total = 0.0, 0, 0

            for inputs, labels in trainloader:
                inputs, labels = inputs.to(device), labels.to(device)
                optimizer.zero_grad()
                outputs = model(inputs)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer.step()
                apply_masks(model, masks) # Keep connectivity fixed

                train_loss += loss.item()
                _, predicted = outputs.max(1)
                train_total += labels.size(0)
                train_correct += predicted.eq(labels).sum().item()
```

```

# --- Evaluation Phase ---
model.eval()
test_loss, test_correct, test_total = 0.0, 0, 0
with torch.no_grad():
    for inputs, labels in testloader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        test_loss += loss.item()
        _, predicted = outputs.max(1)
        test_total += labels.size(0)
        test_correct += predicted.eq(labels).sum().item()

# Record Metrics
history["train_loss"].append(train_loss / len(trainloader))
history["train_acc"].append(100.0 * train_correct / train_total)
history["test_loss"].append(test_loss / len(testloader))
history["test_acc"].append(100.0 * test_correct / test_total)

print(
    f"Epoch {epoch+1}: Train Loss {history['train_loss'][-1]:.4f} | Test Acc {hi
)

# --- Save Joblib ---
experiment_results = {
    "sparsity": s,
    "optimizer": opt_name,
    "train_loss": history["train_loss"],
    "train_acc": history["train_acc"],
    "test_loss": history["test_loss"],
    "test_acc": history["test_acc"],
    "masks": masks, # Required deliverable
    "model_state": model.state_dict(), # For future analysis
}

filename = os.path.join(
    output_dir, f"task_1_{opt_name}_{int(s*100)}_sparsity.joblib"
)
joblib.dump(experiment_results, filename)
print(f"Results saved to {filename}!")

```

Starting Experiment: Sparsity=90.0%, Optimizer=sgd

Epoch 1:	Train Loss	1.4800		Test Acc	46.15%
Epoch 2:	Train Loss	1.0508		Test Acc	57.55%
Epoch 3:	Train Loss	0.8884		Test Acc	54.54%
Epoch 4:	Train Loss	0.8074		Test Acc	51.62%
Epoch 5:	Train Loss	0.7596		Test Acc	69.40%
Epoch 6:	Train Loss	0.7299		Test Acc	64.62%
Epoch 7:	Train Loss	0.7096		Test Acc	67.02%
Epoch 8:	Train Loss	0.6902		Test Acc	68.87%
Epoch 9:	Train Loss	0.6775		Test Acc	70.44%
Epoch 10:	Train Loss	0.6648		Test Acc	69.63%
Epoch 11:	Train Loss	0.6513		Test Acc	70.86%
Epoch 12:	Train Loss	0.6478		Test Acc	60.67%
Epoch 13:	Train Loss	0.6390		Test Acc	70.60%
Epoch 14:	Train Loss	0.6302		Test Acc	71.90%
Epoch 15:	Train Loss	0.6326		Test Acc	70.14%

Results saved to /content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project/results/t

Starting Experiment: Sparsity=90.0%, Optimizer=adam

Epoch 1:	Train Loss	1.5871		Test Acc	48.50%
Epoch 2:	Train Loss	1.1645		Test Acc	58.19%
Epoch 3:	Train Loss	1.0149		Test Acc	60.58%
Epoch 4:	Train Loss	0.9248		Test Acc	66.05%
Epoch 5:	Train Loss	0.8605		Test Acc	64.21%
Epoch 6:	Train Loss	0.8024		Test Acc	67.34%
Epoch 7:	Train Loss	0.7640		Test Acc	68.56%
Epoch 8:	Train Loss	0.7243		Test Acc	70.47%
Epoch 9:	Train Loss	0.6920		Test Acc	70.43%
Epoch 10:	Train Loss	0.6628		Test Acc	70.66%
Epoch 11:	Train Loss	0.6403		Test Acc	71.42%
Epoch 12:	Train Loss	0.6167		Test Acc	71.68%
Epoch 13:	Train Loss	0.5946		Test Acc	73.01%
Epoch 14:	Train Loss	0.5739		Test Acc	72.12%
Epoch 15:	Train Loss	0.5620		Test Acc	73.34%

Results saved to /content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project/results/t

Starting Experiment: Sparsity=95.0%, Optimizer=sgd

Epoch 1:	Train Loss	1.6498		Test Acc	38.35%
Epoch 2:	Train Loss	1.2002		Test Acc	57.12%
Epoch 3:	Train Loss	1.0380		Test Acc	59.63%
Epoch 4:	Train Loss	0.9561		Test Acc	60.00%
Epoch 5:	Train Loss	0.9054		Test Acc	52.98%

Epoch 6: Train Loss 0.8657 | Test Acc 65.12%
Epoch 7: Train Loss 0.8465 | Test Acc 62.64%
Epoch 8: Train Loss 0.8222 | Test Acc 63.60%
Epoch 9: Train Loss 0.8056 | Test Acc 65.55%
Epoch 10: Train Loss 0.7962 | Test Acc 67.65%
Epoch 11: Train Loss 0.7894 | Test Acc 66.20%
Epoch 12: Train Loss 0.7805 | Test Acc 55.36%
Epoch 13: Train Loss 0.7666 | Test Acc 60.81%
Epoch 14: Train Loss 0.7665 | Test Acc 67.60%
Epoch 15: Train Loss 0.7544 | Test Acc 61.46%

Results saved to /content/drive/MyDrive/Posgradados/2025/Duke/ECE 685D/Final project/results/t

Starting Experiment: Sparsity=95.0%, Optimizer=adam

Epoch 1: Train Loss 1.8323 | Test Acc 36.19%
Epoch 2: Train Loss 1.5417 | Test Acc 41.56%
Epoch 3: Train Loss 1.4106 | Test Acc 42.55%
Epoch 4: Train Loss 1.3250 | Test Acc 45.08%
Epoch 5: Train Loss 1.2622 | Test Acc 48.91%
Epoch 6: Train Loss 1.2182 | Test Acc 49.05%
Epoch 7: Train Loss 1.1826 | Test Acc 50.83%
Epoch 8: Train Loss 1.1545 | Test Acc 54.83%
Epoch 9: Train Loss 1.1285 | Test Acc 51.85%
Epoch 10: Train Loss 1.1068 | Test Acc 56.14%
Epoch 11: Train Loss 1.0861 | Test Acc 56.15%
Epoch 12: Train Loss 1.0698 | Test Acc 57.18%
Epoch 13: Train Loss 1.0574 | Test Acc 55.74%
Epoch 14: Train Loss 1.0411 | Test Acc 57.95%
Epoch 15: Train Loss 1.0289 | Test Acc 57.92%

Results saved to /content/drive/MyDrive/Posgradados/2025/Duke/ECE 685D/Final project/results/t

Starting Experiment: Sparsity=98.0%, Optimizer=sgd

Epoch 1: Train Loss 1.7946 | Test Acc 36.59%
Epoch 2: Train Loss 1.5153 | Test Acc 46.35%
Epoch 3: Train Loss 1.4071 | Test Acc 45.13%
Epoch 4: Train Loss 1.3515 | Test Acc 42.58%
Epoch 5: Train Loss 1.3171 | Test Acc 45.01%
Epoch 6: Train Loss 1.2975 | Test Acc 42.05%
Epoch 7: Train Loss 1.2746 | Test Acc 47.86%
Epoch 8: Train Loss 1.2583 | Test Acc 50.42%
Epoch 9: Train Loss 1.2576 | Test Acc 44.14%
Epoch 10: Train Loss 1.2435 | Test Acc 50.08%
Epoch 11: Train Loss 1.2302 | Test Acc 46.97%
Epoch 12: Train Loss 1.2275 | Test Acc 46.45%

```
Epoch 13: Train Loss 1.2237 | Test Acc 44.64%
Epoch 14: Train Loss 1.2136 | Test Acc 36.23%
Epoch 15: Train Loss 1.2156 | Test Acc 50.50%
Results saved to /content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project/results/t
```

Starting Experiment: Sparsity=98.0%, Optimizer=adam

```
Epoch 1: Train Loss 1.9828 | Test Acc 30.40%
Epoch 2: Train Loss 1.7337 | Test Acc 37.47%
Epoch 3: Train Loss 1.6010 | Test Acc 40.73%
Epoch 4: Train Loss 1.5292 | Test Acc 41.95%
Epoch 5: Train Loss 1.4710 | Test Acc 43.88%
Epoch 6: Train Loss 1.4234 | Test Acc 45.62%
Epoch 7: Train Loss 1.3807 | Test Acc 45.25%
Epoch 8: Train Loss 1.3495 | Test Acc 45.60%
Epoch 9: Train Loss 1.3253 | Test Acc 46.64%
Epoch 10: Train Loss 1.3034 | Test Acc 47.59%
Epoch 11: Train Loss 1.2862 | Test Acc 48.19%
Epoch 12: Train Loss 1.2712 | Test Acc 49.35%
Epoch 13: Train Loss 1.2562 | Test Acc 48.64%
Epoch 14: Train Loss 1.2407 | Test Acc 46.10%
Epoch 15: Train Loss 1.2298 | Test Acc 47.28%
Results saved to /content/drive/MyDrive/Posgrados/2025/Duke/ECE 685D/Final project/results/t
```

The printed outputs provide a quick view of training loss and test accuracy across epochs for each configuration. As sparsity increases, performance generally degrades, reflecting the reduced capacity of the network. Differences between optimizers also become more pronounced at higher sparsities, with Adam often converging more smoothly than SGD.

These logs only show a subset of the recorded metrics. Additional information such as training accuracy, test loss, and the full training history are stored in the saved joblib files for more detailed analysis.

Evaluation

Functions

```
def plot_training_and_validation(results):
    sns.set_theme(style="whitegrid", context="talk", palette="deep")

    sparsities = [0.90, 0.95, 0.98]
    opts = ["sgd", "adam"]
```

```

for s in sparsities:
    rows = []

    # --- Reshape data into tidy format ---
    for opt in opts:
        if s in results[opt]:
            n_epochs = len(results[opt][s]["train_loss"])
            for epoch in range(n_epochs):
                rows.append(
                    {
                        "epoch": epoch,
                        "optimizer": opt.upper(),
                        "split": "Train",
                        "loss": results[opt][s]["train_loss"][epoch],
                        "accuracy": results[opt][s]["train_acc"][epoch],
                    }
                )
            rows.append(
                {
                    "epoch": epoch,
                    "optimizer": opt.upper(),
                    "split": "Test",
                    "loss": results[opt][s]["test_loss"][epoch],
                    "accuracy": results[opt][s]["test_acc"][epoch],
                }
            )

    df = pd.DataFrame(rows)

    # --- Create figure ---
    fig, axes = plt.subplots(1, 2, figsize=(16, 6))
    fig.suptitle(
        f"Optimization Analysis at {int(s*100)}% Sparsity",
        fontsize=18,
        weight="bold",
    )

    # --- Loss plot ---
    sns.lineplot(
        data=df,
        x="epoch",
        y="loss",

```

```

        hue="optimizer",
        style="split",
        markers=False,
        dashes={"Train": "", "Test": (4, 2)},
        ax=axes[0],
    )
    axes[0].set_title("Cross-Entropy Loss", fontsize=14)
    axes[0].set_xlabel("Epoch")
    axes[0].set_ylabel("Loss")

    # --- Accuracy plot ---
    sns.lineplot(
        data=df,
        x="epoch",
        y="accuracy",
        hue="optimizer",
        style="split",
        markers=False,
        dashes={"Train": "", "Test": (4, 2)},
        ax=axes[1],
    )
    axes[1].set_title("Accuracy (%)", fontsize=14)
    axes[1].set_xlabel("Epoch")
    axes[1].set_ylabel("Accuracy (%)")

    # --- Final polish ---
    for ax in axes:
        ax.grid(True, alpha=0.3)
        ax.legend(title="")

    plt.savefig(f"img/sparsity_{int(s*100)}.png", dpi=300, bbox_inches="tight")
    plt.tight_layout()
    plt.show()

```

Read results

```
output_dir = "results"
```

```
# Patch torch.load to force CPU
_original_torch_load = torch.load
```

```
def cpu_torch_load(*args, **kwargs):
    kwargs["map_location"] = torch.device("cpu")
    return _original_torch_load(*args, **kwargs)
```

```
torch.load = cpu_torch_load
```

```
results = {}
```

```
for opt_name in optimizers:
    results[opt_name] = {}
    for s in sparsities:
        file_path = os.path.join(
            output_dir, f"task_1_{opt_name}_{int(s*100)}_sparsity.joblib"
        )


        if os.path.exists(file_path):
            data = joblib.load(file_path)
            results[opt_name][s] = data
            print(f"Loaded: {file_path}")
        else:
            print(f"Missing: {file_path}")
```

```
Loaded: results/task_1_sgd_90_sparsity.joblib
Loaded: results/task_1_sgd_95_sparsity.joblib
Loaded: results/task_1_sgd_98_sparsity.joblib
Loaded: results/task_1_adam_90_sparsity.joblib
Loaded: results/task_1_adam_95_sparsity.joblib
Loaded: results/task_1_adam_98_sparsity.joblib
```


Analysis

```
plot_training_and_validation(results)
```

Optimization Analysis at 90% Sparsity



Optimization Analysis at 95% Sparsity



Optimization Analysis at 98% Sparsity

